

# Problem Statement :-

**Q2.** Imagine you have a dataset where you have different features like Age , Gender , Height , Weight , BMI , and Blood Pressure and you have to classify the people into different classes like Normal , Overweight , Obesity , Underweight , and Extreme Obesity by using any 4 different classification algorithms. Now you have to build a model which can classify people into different classes.

Dataset link:- <https://www.kaggle.com/datasets/ankurbajaj9/obesity-levels>

```
In [1]: ## Import the necessary libraries:-
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.preprocessing import LabelEncoder
import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: # Load the dataset
data = pd.read_csv(r"C:\Users\hrush\Downloads\archive (3)\ObesityDataSet_raw_and_data_synthetic.csv")
```

```
In [4]: ## Checking top 5 rows
data.head()
```

```
Out[4]:
```

	Gender	Age	Height	Weight	family_history_with_overweight	FAVC	FCVC	NCP	CAEC	SMOKE	CH2O	SCC	FAF	TUE	CALC
0	Female	21.0	1.62	64.0	yes	no	2.0	3.0	Sometimes	no	2.0	no	0.0	1.0	nc
1	Female	21.0	1.52	56.0	yes	no	3.0	3.0	Sometimes	yes	3.0	yes	3.0	0.0	Sometimes
2	Male	23.0	1.80	77.0	yes	no	2.0	3.0	Sometimes	no	2.0	no	2.0	1.0	Frequently
3	Male	27.0	1.80	87.0	no	no	3.0	3.0	Sometimes	no	2.0	no	2.0	0.0	Frequently
4	Male	22.0	1.78	89.8	no	no	2.0	1.0	Sometimes	no	2.0	no	0.0	0.0	Sometimes

```
In [5]: ## Checking Rows & Columns Availabale in Dataset
data.shape
```

```
Out[5]: (2111, 17)
```

```
In [6]: ## Checking Details Information related with Dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   Gender                                2111 non-null   object
 1   Age                                   2111 non-null   float64
 2   Height                                2111 non-null   float64
 3   Weight                                2111 non-null   float64
 4   family_history_with_overweight        2111 non-null   object
 5   FAVC                                  2111 non-null   object
 6   FCVC                                  2111 non-null   float64
 7   NCP                                   2111 non-null   float64
 8   CAEC                                  2111 non-null   object
 9   SMOKE                                 2111 non-null   object
10   CH2O                                  2111 non-null   float64
11   SCC                                   2111 non-null   object
12   FAF                                   2111 non-null   float64
13   TUE                                   2111 non-null   float64
14   CALC                                  2111 non-null   object
15   MTRANS                                2111 non-null   object
16   NObesidad                             2111 non-null   object
dtypes: float64(8), object(9)
memory usage: 280.5+ KB
```

```
In [7]: ## Checking All Columns name present in dataset
data.columns
```

```
Out[7]: Index(['Gender', 'Age', 'Height', 'Weight', 'family_history_with_overweight',
              'FAVC', 'FCVC', 'NCP', 'CAEC', 'SMOKE', 'CH2O', 'SCC', 'FAF', 'TUE',
              'CALC', 'MTRANS', 'NObesidad'],
              dtype='object')
```

```
In [8]: ## Checking Statistical Analysis of Dataset
data.describe()
```

Out[8]:

	Age	Height	Weight	FCVC	NCP	CH2O	FAF	TUE
count	2111.000000	2111.000000	2111.000000	2111.000000	2111.000000	2111.000000	2111.000000	2111.000000
mean	24.312600	1.701677	86.586058	2.419043	2.685628	2.008011	1.010298	0.657866
std	6.345968	0.093305	26.191172	0.533927	0.778039	0.612953	0.850592	0.608927
min	14.000000	1.450000	39.000000	1.000000	1.000000	1.000000	0.000000	0.000000
25%	19.947192	1.630000	65.473343	2.000000	2.658738	1.584812	0.124505	0.000000
50%	22.777890	1.700499	83.000000	2.385502	3.000000	2.000000	1.000000	0.625350
75%	26.000000	1.768464	107.430682	3.000000	3.000000	2.477420	1.666678	1.000000
max	61.000000	1.980000	173.000000	3.000000	4.000000	3.000000	3.000000	2.000000

```
In [9]: ## Checking Information Related with Dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Gender                                     2111 non-null   object
1   Age                                       2111 non-null   float64
2   Height                                   2111 non-null   float64
3   Weight                                   2111 non-null   float64
4   family_history_with_overweight          2111 non-null   object
5   FAVC                                     2111 non-null   object
6   FCVC                                     2111 non-null   float64
7   NCP                                       2111 non-null   float64
8   CAEC                                     2111 non-null   object
9   SMOKE                                    2111 non-null   object
10  CH2O                                     2111 non-null   float64
11  SCC                                       2111 non-null   object
12  FAF                                       2111 non-null   float64
13  TUE                                       2111 non-null   float64
14  CALC                                     2111 non-null   object
15  MTRANS                                    2111 non-null   object
16  NObeyesdad                               2111 non-null   object
dtypes: float64(8), object(9)
memory usage: 280.5+ KB
```

```
In [10]: ## Checking All Columns Available in dataset
data.columns
```

```
Out[10]: Index(['Gender', 'Age', 'Height', 'Weight', 'family_history_with_overweight',
               'FAVC', 'FCVC', 'NCP', 'CAEC', 'SMOKE', 'CH2O', 'SCC', 'FAF', 'TUE',
               'CALC', 'MTRANS', 'NObeyesdad'],
              dtype='object')
```

```
In [11]: # Preprocess the dataset
encoder = LabelEncoder()
data['Gender'] = encoder.fit_transform(data['Gender'])
data['family_history_with_overweight'] = encoder.fit_transform(data['family_history_with_overweight'])
data['FAVC'] = encoder.fit_transform(data['FAVC'])
data['CAEC'] = encoder.fit_transform(data['CAEC'])
data['SMOKE'] = encoder.fit_transform(data['SMOKE'])
data['SCC'] = encoder.fit_transform(data['SCC'])
data['CALC'] = encoder.fit_transform(data['CALC'])
data['MTRANS'] = encoder.fit_transform(data['MTRANS'])
data['NObeyesdad'] = encoder.fit_transform(data['NObeyesdad'])
```

```
In [12]: ## Checking Details Information related with Dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Gender                                     2111 non-null   int32
1   Age                                       2111 non-null   float64
2   Height                                   2111 non-null   float64
3   Weight                                   2111 non-null   float64
4   family_history_with_overweight          2111 non-null   int32
5   FAVC                                     2111 non-null   int32
6   FCVC                                     2111 non-null   float64
7   NCP                                       2111 non-null   float64
8   CAEC                                     2111 non-null   int32
9   SMOKE                                    2111 non-null   int32
10  CH20                                     2111 non-null   float64
11  SCC                                       2111 non-null   int32
12  FAF                                       2111 non-null   float64
13  TUE                                       2111 non-null   float64
14  CALC                                     2111 non-null   int32
15  MTRANS                                    2111 non-null   int32
16  NObeyesdad                              2111 non-null   int32
dtypes: float64(8), int32(9)
memory usage: 206.3 KB
```

```
In [13]: # Split the dataset into features (X) and target (y)
X = data.drop('NObeyesdad', axis=1)
y = data['NObeyesdad']
```

```
In [14]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [15]: # Decision Tree Classifier
dt_clf = DecisionTreeClassifier()
dt_clf.fit(X_train, y_train)
dt_predictions = dt_clf.predict(X_test)
```

```
In [16]: # Logistic Regression Classifier
lr_clf = LogisticRegression()
lr_clf.fit(X_train, y_train)
lr_predictions = lr_clf.predict(X_test)
```

```
In [17]: # Random Forest Classifier
rf_clf = RandomForestClassifier()
rf_clf.fit(X_train, y_train)
rf_predictions = rf_clf.predict(X_test)
```

```
In [18]: # Support Vector Machine (SVM) Classifier
svm_clf = SVC()
svm_clf.fit(X_train, y_train)
svm_predictions = svm_clf.predict(X_test)
```

```
In [19]: # Print classification reports for each classifier
print("Decision Tree Classifier:")
print(classification_report(y_test, dt_predictions))
```

```
Decision Tree Classifier:
              precision    recall  f1-score   support

     0       0.90      0.98      0.94         56
     1       0.88      0.84      0.86         62
     2       0.97      0.92      0.95         78
     3       0.95      0.95      0.95         58
     4       1.00      1.00      1.00         63
     5       0.89      0.91      0.90         56
     6       0.94      0.96      0.95         50

 accuracy          0.94          0.94          0.94         423
 macro avg         0.93          0.94          0.94         423
 weighted avg      0.94          0.94          0.94         423
```

```
In [20]: print("Logistic Regression Classifier:")
print(classification_report(y_test, lr_predictions))
```

Logistic Regression Classifier:				
	precision	recall	f1-score	support
0	0.74	0.93	0.83	56
1	0.53	0.42	0.47	62
2	0.58	0.60	0.59	78
3	0.82	0.84	0.83	58
4	0.90	1.00	0.95	63
5	0.54	0.38	0.44	56
6	0.35	0.38	0.37	50
accuracy			0.65	423
macro avg			0.64	423
weighted avg			0.64	423

```
In [21]: print("Random Forest Classifier:")
print(classification_report(y_test, rf_predictions))
```

Random Forest Classifier:				
	precision	recall	f1-score	support
0	1.00	0.96	0.98	56
1	0.89	0.92	0.90	62
2	0.99	0.95	0.97	78
3	0.97	0.98	0.97	58
4	1.00	1.00	1.00	63
5	0.88	0.89	0.88	56
6	0.92	0.94	0.93	50
accuracy			0.95	423
macro avg			0.95	423
weighted avg			0.95	423

```
In [22]: print("SVM Classifier:")
print(classification_report(y_test, svm_predictions))
```

SVM Classifier:				
	precision	recall	f1-score	support
0	0.71	0.88	0.78	56
1	0.48	0.34	0.40	62
2	0.65	0.33	0.44	78
3	0.77	0.41	0.54	58
4	0.56	1.00	0.72	63
5	0.47	0.48	0.47	56
6	0.43	0.58	0.49	50
accuracy			0.57	423
macro avg			0.58	423
weighted avg			0.59	423

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js